

Introduction

Concept learning can be formulated as a problem of searching through a predefined space of potential hypotheses for the hypothesis that best fits the training examples

Other definition . Inferring a boolean-valued function from training examples of its input and output.

Hypothesis

indicate by a "?" that any value is acceptable for this attribute,
specify a single required value (e.g., Warm) for the attribute, or
indicate by a "0" that no value is acceptable.

$\langle ?, \text{Cold}, \text{High}, ?, ?, ? \rangle$

Example	<i>Sky</i>	<i>AirTemp</i>	<i>Humidity</i>	<i>Wind</i>	<i>Water</i>	<i>Forecast</i>	<i>EnjoySport</i>
1	Sunny	Warm	Normal	Strong	Warm	Same	Yes
2	Sunny	Warm	High	Strong	Warm	Same	Yes
3	Rainy	Cold	High	Strong	Warm	Change	No
4	Sunny	Warm	High	Strong	Cool	Change	Yes

Learning Task

The most **general hypothesis-that** every day is a positive example-is represented by $(?, ?, ?, ?, ?, ?)$

and the most **specific** possible hypothesis-that no day is a positive example-is represented by $(0,0,0,0,0,0)$

To summarize, the **EnjoySport concept learning task** requires learning the set of days for which EnjoySport = **yes**, describing this set by a conjunction of constraints over the instance attributes.

- **Given:**
 - Instances X : Possible days, each described by the attributes
 - *Sky* (with possible values *Sunny*, *Cloudy*, and *Rainy*),
 - *AirTemp* (with values *Warm* and *Cold*),
 - *Humidity* (with values *Normal* and *High*),
 - *Wind* (with values *Strong* and *Weak*),
 - *Water* (with values *Warm* and *Cool*), and
 - *Forecast* (with values *Same* and *Change*).
 - Hypotheses H : Each hypothesis is described by a conjunction of constraints on the attributes *Sky*, *AirTemp*, *Humidity*, *Wind*, *Water*, and *Forecast*. The constraints may be “?” (any value is acceptable), “ \emptyset ” (no value is acceptable), or a specific value.
 - Target concept c : $EnjoySport : X \rightarrow \{0, 1\}$
 - Training examples D : Positive and negative examples of the target function (see Table 2.1).
- **Determine:**
 - A hypothesis h in H such that $h(x) = c(x)$ for all x in X .

Inductive learning

When learning the target concept, the learner is presented a set of training examples, each consisting of an instance x from X , along with its target concept value $c(x)$

Instances for which $c(x) = 1$ are called **positive examples**, or members of the target concept. Instances for which $c(x) = 0$ are called **negative examples**, or nonmembers

We use the symbol H to denote the set of all possible hypotheses that the learner may consider regarding the identity of the target concept


The goal of the learner is to find a hypothesis h such that
 $h(x) = c(x)$ for all x in X .

The inductive learning hypothesis.

Any **hypothesis** found to approximate the target function well **over a sufficiently large set** of training examples will also approximate the target function well over other **unobserved examples**.

Concept Learning as a search

Concept learning can be viewed as the task of searching through a large space of hypotheses implicitly defined by the hypothesis representation



General-to-Specific Ordering

Many algorithms for concept learning organize the search through the hypothesis space by relying on a very useful structure that exists for any concept Learning problem: a general-to-specific ordering of hypotheses.

Definition: Let h_j and h_k be boolean-valued functions defined over X . Then h_j is **more_general_than_or_equal_to** h_k (written $h_j \succeq_g h_k$) if and only if

$$(\forall x \in X)[(h_k(x) = 1) \rightarrow (h_j(x) = 1)]$$

Given hypotheses h_j and h_k , h_j is more-general-than-or-equal-to h_k if and only if any instance that satisfies h_k also satisfies h_j .

$$h_1 = \langle \text{Sunny}, ?, ?, \text{Strong}, ?, ? \rangle$$

$$h_2 = \langle \text{Sunny}, ?, ?, ?, ?, ? \rangle$$

Find-S

One way is to begin with the most specific possible hypothesis in H , then generalize this hypothesis each step, the hypothesis is generalized only as far as necessary to cover the new positive example


1. Initialize h to the most specific hypothesis in H
 2. For each positive training instance x
 - For each attribute constraint a_i in h
 - If the constraint a_i is satisfied by x
Then do nothing
 - Else replace a_i in h by the next more general constraint that is satisfied by x
 3. Output hypothesis h
-

$h \leftarrow \langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle$

$h \leftarrow \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, \text{Warm}, \text{Same} \rangle$

$h \leftarrow \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, ?, ? \rangle$

Find-S



VERSION SPACES AND THE CANDIDATE-ELIMINATION ALGORITHM

TARGET: output a description of the set of all hypotheses consistent with the training example

Definition: A hypothesis h is **consistent** with a set of training examples D if and only if $h(x) = c(x)$ for each example $\langle x, c(x) \rangle$ in D .

$$\text{Consistent}(h, D) \equiv (\forall \langle x, c(x) \rangle \in D) h(x) = c(x)$$

Definition: The **version space**, denoted $VS_{H,D}$, with respect to hypothesis space H and training examples D , is the subset of hypotheses from H consistent with the training examples in D .

$$VS_{H,D} \equiv \{h \in H \mid \text{Consistent}(h, D)\}$$

LIST-THEN-ELIMINATE ALGORITHM


The LIST-THEN-ELIMINATE algorithm first initializes the version space to contain all hypotheses in H , then eliminates any hypothesis found inconsistent with any training example.

The LIST-THEN-ELIMINATE Algorithm

1. $VersionSpace \leftarrow$ a list containing every hypothesis in H
 2. For each training example, $\langle x, c(x) \rangle$
 remove from $VersionSpace$ any hypothesis h for which $h(x) \neq c(x)$
 3. Output the list of hypotheses in $VersionSpace$
-

Unfortunately, it requires exhaustively enumerating all hypotheses in H —an Unrealistic requirement for all but the most trivial hypothesis spaces.

A more compact representation for version spaces



The version space includes all six hypotheses shown here, but can be represented more simply by S and G. Arrows indicate instances of the relation *more-general-than*

A more compact representation

Output for Find_S

$$h = \langle \text{Sunny}, \text{Warm}, ?, \text{Strong}, ?, ? \rangle$$

this is just one of six different hypotheses from H that are consistent with these training examples

Definition: The **general boundary** G , with respect to hypothesis space H and training data D , is the set of maximally general members of H consistent with D .

$$G \equiv \{g \in H \mid \text{Consistent}(g, D) \wedge (\neg \exists g' \in H)[(g' >_g g) \wedge \text{Consistent}(g', D)]\}$$

Definition: The **specific boundary** S , with respect to hypothesis space H and training data D , is the set of minimally general (i.e., maximally specific) members of H consistent with D .

$$S \equiv \{s \in H \mid \text{Consistent}(s, D) \wedge (\neg \exists s' \in H)[(s >_g s') \wedge \text{Consistent}(s', D)]\}$$

A more compact representation

Output for Find_S

$$h = \langle \textit{Sunny}, \textit{Warm}, ?, \textit{Strong}, ?, ? \rangle$$

this is just one of six different hypotheses from H that are consistent with these training examples

Definition: The **general boundary** G , with respect to hypothesis space H and training data D , is the set of maximally general members of H consistent with D .

$$G \equiv \{g \in H \mid \textit{Consistent}(g, D) \wedge (\neg \exists g' \in H)[(g' >_g g) \wedge \textit{Consistent}(g', D)]\}$$

Definition: The **specific boundary** S , with respect to hypothesis space H and training data D , is the set of minimally general (i.e., maximally specific) members of H consistent with D .

$$S \equiv \{s \in H \mid \textit{Consistent}(s, D) \wedge (\neg \exists s' \in H)[(s >_g s') \wedge \textit{Consistent}(s', D)]\}$$

the version space is precisely the set of hypotheses contained in G , plus those contained in S , plus those that lie between G and S in the partially ordered hypothesis space

CANDIDATE-ELIMINATION LEARNING ALGORITHM

Computes the version space containing all hypotheses from H that are consistent with an observed sequence of training examples.

1- Initializing the **G** boundary set to contain the most general hypothesis in H

2- initializing the **S** boundary set to contain the most specific (least general) hypothesis

$$G_0 \leftarrow \{ \langle ?, ?, ?, ?, ?, ? \rangle \} \quad S_0 \leftarrow \{ \langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle \}$$

Initialize G to the set of maximally general hypotheses in H

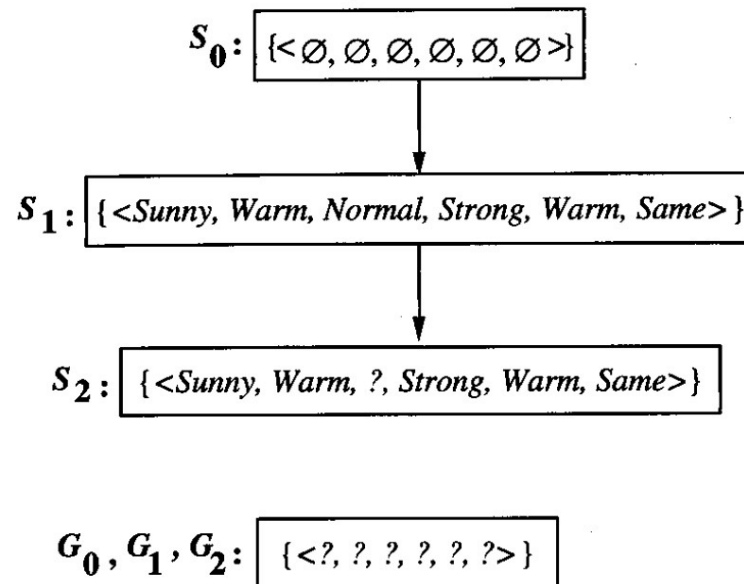
Initialize S to the set of maximally specific hypotheses in H

For each training example d , do

- If d is a positive example
 - Remove from G any hypothesis inconsistent with d
 - For each hypothesis s in S that is not consistent with d
 - Remove s from S
 - Add to S all minimal generalizations h of s such that
 - h is consistent with d , and some member of G is more general than h
 - Remove from S any hypothesis that is more general than another hypothesis in S
- If d is a negative example
 - Remove from S any hypothesis inconsistent with d
 - For each hypothesis g in G that is not consistent with d
 - Remove g from G
 - Add to G all minimal specializations h of g such that
 - h is consistent with d , and some member of S is more specific than h
 - Remove from G any hypothesis that is less general than another hypothesis in G

CANDIDATE-ELIMINATION LEARNING ALGORITHM

Step 1: Training examples 1 and 2 force the boundary to become S more general, as in the FIND-S algorithm. They have no effect on the boundary G



Training examples:


1. $\langle \text{Sunny}, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Enjoy Sport} = \text{Yes}$
2. $\langle \text{Sunny}, \text{Warm}, \text{High}, \text{Strong}, \text{Warm}, \text{Same} \rangle, \text{Enjoy Sport} = \text{Yes}$

positive training examples may force the S boundary of the version space to become increasingly general.
Negative training examples play the complementary role of forcing the G boundary to become increasingly specific

CANDIDATE-ELIMINATION LEARNING ALGORITHM

Step 2: Training example 3 is a negative example that forces the G_2 boundary to be specialized to G_3

$S_2, S_3: \{ \langle \text{Sunny, Warm, ?, Strong, Warm, Same} \rangle \}$



Training Example:

3. $\langle \text{Rainy, Cold, High, Strong, Warm, Change} \rangle, \text{EnjoySport}=\text{No}$

positive training examples may force the S boundary of the version space to become increasingly general.
Negative training examples play the complementary role of forcing the G boundary to become increasingly specific